

```
#1
#Dataset downloaded
#Uploading the dataset
```

```
from google.colab import files
uploaded = files.upload()
```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving abalone.csv to abalone (1).csv

```
#2
#Loading the dataset
```

```
import pandas as pd
import numpy as np
```

```
df = pd.read_csv(r'abalone.csv')
df.head(3)
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.15	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.07	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.21	9

```
#3 Univariate analysis
```

```
import matplotlib.pyplot as plt
```


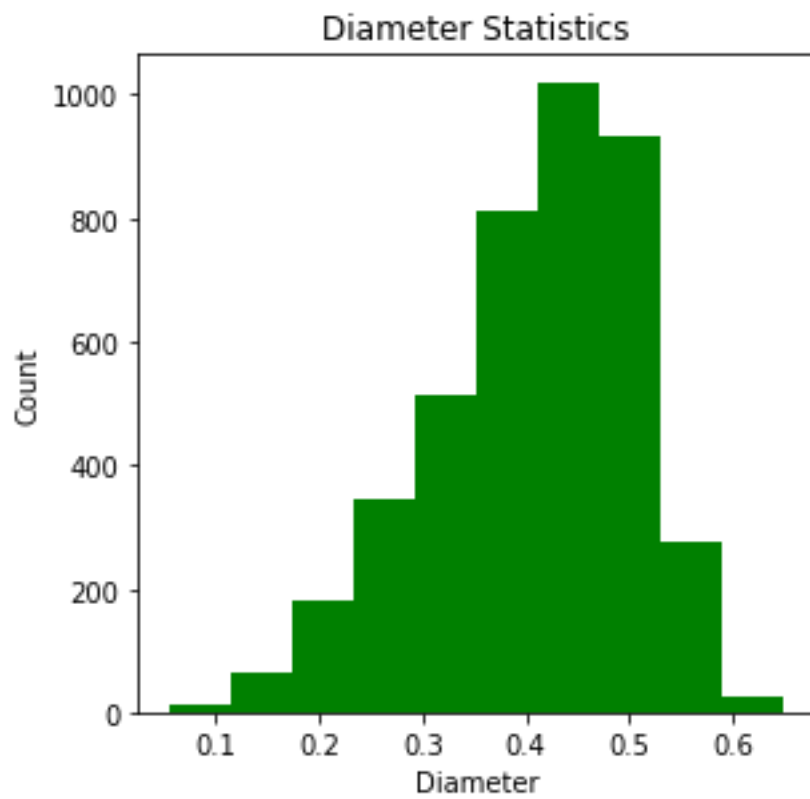


```
plt.figure(figsize=(10,10))
counts = df['Sex'].value_counts()
plt.subplot(221)
plt.pie(counts, labels = counts.index)
plt.title('Sex Statistics')
plt.xlabel('Sex')
plt.ylabel('Count')
```

```
plt.subplot(222)
plt.hist(df['Height'],color='blue')
plt.title('Height Statistics')
plt.xlabel('Height')
plt.ylabel('Count')
```

```
plt.subplot(223)
plt.hist(df['Diameter'],color='green')
plt.title('Diameter Statistics')
plt.xlabel('Diameter')
plt.ylabel('Count')
```

```
plt.subplot(224)
df['Whole weight'].plot(kind='density')
plt.title('Whole weight Statistics')
```

```
plt.xlabel('Whole weight')
plt.ylabel('Count')
Text(0, 0.5, 'Count')
```



```
#3 Bivariate Analysis
import seaborn as sns
```

```
plt.figure(figsize=(15,15))
```

```
#Categorical vs Categorical
```

```
plt.subplot(2,1,1)
```

```
sns.countplot(data = df, x = 'Rings', hue = 'Sex')
```

```
plt.figure(figsize=(15,15))
```

```
#Continuous vs Continuous
```


```
plt.subplot(2,2,1)
```

```
sns.stripplot(x='Rings', y='Viscera weight', data=df, palette='rainbow')
```

```
plt.subplot(2,2,2)
```


```
sns.stripplot(x='Rings', y='Height', data=df, palette='rainbow')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f68cee0d90>
```





```
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cd0dec90>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cd0a21d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cd0586d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cd010bd0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccfd2110>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cecc1f10>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cd3c71d0>],  
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cece1710>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cf0a1c90>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cecfdded0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cef23d10>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cf9010d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cefd8d90>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccf943d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccf4d9d0>],  
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccf0f050>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccce5610>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cce7ec10>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cce41250>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccdf7850>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccdafa50>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccd71490>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccd28a90>],  
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cccec0d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccd246d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cccd8cd0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccc9d310>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccc53990>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccc07f90>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccbc95d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccb80bd0>],  
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccb46210>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccafa810>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68ccab1e10>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cca7b050>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cca2f510>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc9e6a10>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cca1cf10>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc9e0450>],  
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc997950>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc94de50>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc912390>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc8c8890>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc880d90>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc8442d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc7f97d0>,  
<matplotlib.axes._subplots.AxesSubplot object at 0x7f68cc7aecd0>]],  
dtype=object)
```



#4

df.describe()

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
<b>max</b>	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4177 entries, 0 to 4176

Data columns (total 9 columns):

# Column Non-Null Count Dtype

```
---
0 Sex 4177 non-null object
1 Length 4177 non-null float64
2 Diameter 4177 non-null float64
3 Height 4177 non-null float64
4 Whole weight 4177 non-null float64
5 Shucked weight 4177 non-null float64
6 Viscera weight 4177 non-null float64
7 Shell weight 4177 non-null float64
8 Rings 4177 non-null int64
```

dtypes: float64(7), int64(1), object(1)

memory usage: 293.8+ KB

#5 Handling missing values

df.isnull().sum()

```
Sex      0
Length   0
Diameter  0
Height   0
Whole weight  0
Shucked weight  0
```

```

Viscera weight    0
Shell weight      0
Rings              0
dtype: int64
#For Continuous variables
df['Length'].fillna(df['Length'].mean(), inplace=True)
df['Height'].fillna(df['Height'].mean(), inplace=True)
df['Diameter'].fillna(df['Diameter'].mean(), inplace=True)
df['Whole weight'].fillna(df['Whole weight'].mean(), inplace=True)
df['Shucked weight'].fillna(df['Shucked weight'].mean(), inplace=True)
df['Viscera weight'].fillna(df['Viscera weight'].mean(), inplace=True)
df['Shell weight'].fillna(df['Shell weight'].mean(), inplace=True)
df['Rings'].fillna(df['Rings'].mean(), inplace=True)

#For Categorical variables
df['Sex'].fillna(df['Sex'].mode(), inplace=True)

#Ensuring again
print(df.isnull().sum())


print("\nSex : ', df['Sex'].unique())

Sex          0
Length       0
Diameter     0
Height       0
Whole weight  0
Shucked weight  0
Viscera weight  0
Shell weight  0
Rings        0
dtype: int64

\Sex : ['M' 'F' 'I']
#6 Outlier detection - box plot
import seaborn as sns
sns.boxplot(df['Length'])
print('No. of Outliers : ', (df['Length'] < 0.2).sum())
No. of Outliers : 43

```





*#6 Outlier detection - zscore*

```
from scipy import stats
```

```
zscore = np.abs(stats.zscore(df['Diameter']))
```

```
print(zscore)
```

```
print('No. of Outliers : ', np.shape(np.where(zscore>3)))
```

```
0    0.432149
```

```
1    1.439929
```

```
2    0.122130
```

```
3    0.432149
```

```
4    1.540707
```

```
...
```

```
4172  0.424464
```

```
4173  0.323686
```

```
4174  0.676409
```

```
4175  0.777187
```

```
4176  1.482634
```

```
Name: Diameter, Length: 4177, dtype: float64
```


```
No. of Outliers : (1, 13)
```

*#6 Outlier detection - box plot*

```
sns.boxplot(df['Whole weight'])
```

```
print('No. of Outliers : ', (df['Whole weight'] > 2.2).sum())
```

```
No. of Outliers : 33
```



*#6 Outlier detection - Scatter plot*

**import** matplotlib.pyplot **as** plt

fig, ax = plt.subplots(figsize = (5,3))

ax.scatter(df['Height'], df['Rings'])


*# x-axis label*

ax.set\_xlabel('Height')

*# y-axis label*

ax.set\_ylabel('Rings')

plt.show()



*#6 Outlier detection - IQR*

```
Q1 = df['Shucked weight'].quantile(0.25)
Q3 = df['Shucked weight'].quantile(0.75)
IQR = Q3 - Q1
print(IQR)
upper=Q3 + 1.5 * IQR
lower=Q1 - 1.5 * IQR
count = np.size(np.where(df['Shucked weight'] > upper))
count = count + np.size(np.where(df['Shucked weight'] < lower))
print('No. of outliers : ', count)
```

0.316

No. of outliers : 48

0.316

No. of outliers : 48

*#6 Outlier detection - 3 sigma*

```
upper = df['Viscera weight'].mean() + (3 * df['Viscera weight'].std())
lower = df['Viscera weight'].mean() - (3 * df['Viscera weight'].std())
columns = df[ ( df['Viscera weight'] > upper ) | ( df['Viscera weight'] < lower ) ]
print('Upper range : ', upper)
print('Lower range : ', lower)
print('No. of Outliers : ', len(columns))
```

Upper range : 0.5094363586315791

Lower range : -0.1482491429265277

No. of Outliers : 22

*#6 Outlier detection - 3 sigma*

```
upper = df['Shell weight'].mean() + (3 * df['Shell weight'].std())
lower = df['Shell weight'].mean() - (3 * df['Shell weight'].std())
columns = df[ ( df['Shell weight'] > upper ) | ( df['Shell weight'] < lower ) ]
print('Upper range : ', upper)
print('Lower range : ', lower)
print('No. of Outliers : ', len(columns))
```

Upper range : 0.6564388680356763

Lower range : -0.17877714909864026

No. of Outliers : 27

*#6 Removing Outliers*

```
columns = ['Diameter', 'Length', 'Whole weight', 'Height', 'Shucked weight', 'Viscera weight', 'Shell weight']
```

**for i in columns:**

```
    Q1=df[i].quantile(0.25)
    Q3=df[i].quantile(0.75)
    IQR=Q3-Q1
    upper=Q3+1.5*IQR
    lower=Q1-1.5*IQR
    df[i]=np.where(df[i] > upper, upper, df[i])
    df[i]=np.where(df[i] < lower, lower, df[i])
```

*#After outlier removal*

```
columns = ['Diameter', 'Length', 'Whole weight', 'Height', 'Shucked weight', 'Viscera weight', 'Shell weight']
```

**for i in columns:**

```
    Q1 = df[i].quantile(0.25)
```

```
    Q3 = df[i].quantile(0.75)
```

```

IQR = Q3 - Q1
upper=Q3 + 1.5 * IQR
lower=Q1 - 1.5 * IQR
count = np.size(np.where(df[i] >upper))
count = count + np.size(np.where(df[i] <lower))
print('No. of outliers in ', i, ' : ', count)

```

```

No. of outliers in Diameter : 0
No. of outliers in Length : 0
No. of outliers in Whole weight : 0
No. of outliers in Height : 0
No. of outliers in Shucked weight : 0
No. of outliers in Viscera weight : 0
No. of outliers in Shell weight : 0

```

*#7 Label Encoding*

```

print('Before encoding : ', df['Sex'][0])
df['Sex'] = df['Sex'].astype('category')

```

```

df['Sex'] = df['Sex'].cat.codes
print('After encoding : ', df['Sex'][0])

```

```

df['Sex'].dtype

```

Before encoding : M

After encoding : 2

```

dtype('int8')

```

```

df.dtypes

```

```

Sex          int8
Length       float64
Diameter     float64
Height       float64
Whole weight float64
Shucked weight float64
Viscera weight float64
Shell weight float64
Rings        int64
dtype: object

```

*#Changing target column(Exited) as the last column*

```

Exit = df['Rings']
df = df.drop('Rings', axis=1)
df = df.join(Exit)

```

```

df.head(1)

```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.514	0.2245	0.101	0.15	15

*#8 Splitting dependent and independent variables*

```

X = df.iloc[:, :-1]
print('Independent : \n', X.head(2))

```

```

Y = df.iloc[:, -1]
print('Dependent : \n', Y.head(2))

```

Independent :

	Sex	Length	Diameter	Height	Whole weight	Shucked weight \
0	2	0.455	0.365	0.095	0.5140	0.2245
1	2	0.350	0.265	0.090	0.2255	0.0995

	Viscera weight	Shell weight
0	0.1010	0.15
1	0.0485	0.07

Dependent :

0 15

1 7

Name: Rings, dtype: int64

*#9 Scaling*

```
from sklearn.preprocessing import RobustScaler
scaler=RobustScaler()
```

```
print('Before scaling : \n', X[1:3])
```

```
X = scaler.fit_transform(X)
```

```
print("\n\nAfter scaling : \n', X[1:3])
```

Before scaling :

	Sex	Length	Diameter	Height	Whole weight	Shucked weight \
1	2	0.35	0.265	0.090	0.2255	0.0995
2	0	0.53	0.420	0.135	0.6770	0.2565

	Viscera weight	Shell weight
1	0.0485	0.07
2	0.1415	0.21

After scaling :

```
[[ 0.5      -1.18181818 -1.23076923 -1.      -0.80674631 -0.74841772
  -0.76802508 -0.8241206 ]
 [-0.5      -0.09090909 -0.03846154 -0.1      -0.17217147 -0.25158228
  -0.18495298 -0.12060302]]
```

*#10 Splitting data into training and test set*

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size = 0.2)
```

```
print(X_train.shape, X_test.shape, Y_train.shape, Y_test.shape)
```

```
(3341, 8) (836, 8) (3341,) (836,)
```

*#11 Build the model*

```
from sklearn import linear_model as lin_mod
model = lin_mod.LinearRegression()
```

*#12 Training the model*

```
LinearRegression = model.fit(X_train, Y_train)
```

```
r_sq = model.score(X_train, Y_train)
```

```
print(f"Determination coeeficient: {r_sq}")
```

Determination coeeficient: 0.5394019894190405

*#13 Testing the model*

```
Y_pred = model.predict(X_test)
```

```
print('Predicted values : ', Y_pred[:5])
```

```
Predicted values : [13.18842068 14.21223424  8.63255896  5.61492745  9.05210334]
```

*#Metrics*

```
from sklearn.metrics import r2_score
```

```
from sklearn.metrics import mean_squared_error
```

```
r2 = r2_score(Y_test, Y_pred)
mean_squared_error = mean_squared_error(Y_test, Y_pred)
rmse = (np.sqrt(mean_squared_error))

print('R2 score : ', r2)
print('Mean squared error : ', mean_squared_error)
print('Root Mean squared error : ', rmse)

R2 score : 0.5244781220124035
Mean squared error : 4.597226769271034
Root Mean squared error : 2.144114448734263
```